

Finding a balance between reinforcement and evolution

Filippo Balzarini¹, Jason Kaxiras¹, Melvin Gode¹

^a*Department of Computer Science, Uppsala University, Uppsala, Sweden*

May, 2024

Abstract

$$R = (3.05 \pm 0.02) \cdot 10^{-6} \text{ kg/s}^2$$

1. Introduction

In the field of Machine learning, several methods can be used to solve the same problem. For example, trying to find the local minima using evolutionary algorithms or supervised learning in the form of backpropagation. Both come with advantages and disadvantages depending on if we are after precision or just want to find the local minima fast, but also depending on the nature of our problem.

More specifically interesting is the comparison of the performance of Genetic Algorithms (GA) and Reinforcement Learning (RL) techniques in the context of a game environment. On one hand, in reinforcement learning the agent engages a dynamic and evolving environment by taking actions that affect it to accomplish a specific job. On the other hand, we have evolutionary algorithms that employ evolutionary principles for automated and concurrent problem-solving by drawing inspiration from populations of interacting organisms. Despite their apparent dissimilarities, RL and GA both tackle the same fundamental issue: optimizing a function. This entails maximizing an agent's reward in RL and the fitness function in evolutionary algorithms, respectively, particularly in environments where the parameters may be unknown [drugan2019reinforcement].

This paper focuses on comparing Reinforcement learning and Genetic Algorithms by having them balance a cartpole in 500 moves. More specifically it is a problem in nonlinear dynamics where an inverted pendulum is balancing in a cart. The aim or final goal of both RL and GA are to keep it the system balanced until they run out of moves. the environment will be described in more detail under the environment part.




Figure 1: the cartpole in 2D graphics

Other related research on this topic includes for [drugan2019reinforcement] which focuses on a comprehensive overview of recent trends in the field rather than comparisons of subclasses of algorithms or particular aspects of RL and GA. Several works focus on combining these two methods for machine learning by either using GA to train RL or vice versa such as [eiben2007reinforcement] where the authors try to use Reinforcement learning to tune the parameters of GA. Papers such as [khadka2018evolution] explore the opposite combination of training RL using GA. It is important to mention that the implementation of the reinforcement learning algorithm that is used is based on the work of *JackFurby* [JackFurbyCartPole].

2. Background

2.1. Reinforcement Learning

Reinforcement learning is defined as the problem that an agent tries to solve by learning behaviour through trial and error with its environment. In other words programming an agent through rewards and punishments rather than how to specifically solve the task itself[kaelbling1996reinforcement] as depicted in figure ??.




Figure 2: Graph representing reinforcement learning

The first concept crucial for reinforcement learning is the *reward function* which is objective feedback from the environment. It is usually scalar values that are associated with state-action pairs. High rewards are usually associated by state-action pairs which are beneficial for the agent to be situated in, whereas negative rewards would then be disadvantageous states or *hazardous* for the agent to be in. Essentially, what is good and bad for the agent in the environment. The sole objective of the agent is then to maximize this reward [sutton1999reinforcement].

Naturally, we have to define *state* and *action*, which compared to the rest of the concepts have a very general definition. That being the latter is a decision of some sort and the former a factor that has to be taken into consideration when taking an action.

2.1.1. Temporal difference learning

A central class of methods in reinforcement learning is *temporal difference learning*. It refers to a class of methods in which the learning is based on the difference between temporally successive predictions. It aims to adjust the learner's current expectation for the present input pattern so that it more accurately aligns with the subsequent prediction at the following time step. Unlike Monte Carlo methods and other methods in temporal learning, it updates its estimated value function at every step. [tesauro1995temporal].

In temporal learning, there are several submethods or rather algorithms such as SARSA, Q-learning, TD-Lambda and more [eiben2007reinforcement].

2.1.2. Q learning

Q-learning is an algorithm where the environment can be constituted by a controlled Markov process where the agent is controlling it [watkins1992q]. The agent chooses an action and accordingly gets

rewarded for it. Q-learning uses the Markov chains to calculate the max reward that can be accumulated by the next state-action and updates towards that as shown in the equation below.

$$Q(s, a) := Q(s, a) + \eta[r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \quad (1)$$

Equation ?? is the *value* or *update* equation which is responsible for mapping the different states based on their estimated long-term reward in Q-learning. Here $Q(s, a)$ is the current state of the agent, r is the reward, η is the learning rate, and $Q(s', a')$ is the next state. An important variable here is γ which represents the discount factor. This is used to limit the Markov chain to a limited finite number so they don't end up infinite. This controls how many steps into the future the agent will try to estimate.

2.2. Genetic Algorithms

Genetic algorithms are computational models based on the concept of evolution as seen in biology. Similarly to how organisms evolve by natural selection and sexual reproduction, programs can also simulate these processes and behave in a similar fashion to organisms in order to solve a specific problem. In a general sense, natural selection is the process which determines which individuals get to survive by some test of fitness. After the best-fitted are selected, the creation or reproduction of the next generation starts. Reproduction is then the method in which the mixing of genes in the remaining population happens and gets passed to the offspring [holland1992genetic].




Figure 3: visual representation of genetic algorithms

By starting with a population of individuals which are created randomly, we have an initial population with variation amongst the individuals. The DNA, which is essentially the code of the gene, can be represented by a string of bits. These strings of bits can be thought of as potential solutions to the problem. Due to the variation in the population, some individuals will be better *fit*, which then will be selected to remain. In the final stage, the remaining individuals will mix their bit strings to produce individuals for the next generation. These steps will be

continually done for some number of generations [forrest1996genetic].

It is important however that we reduce the genetic drift and keep track of the best solutions that have been produced by the previous generation. To do that we employ a method called Elitism. In elitism compared with traditional reproduction the most fitting individual are copied to the next generation without any alteration. In that way the best solution of each generation is always preserved and adds selective pressure and improve convergence speed [du2018elitism].

3. Method

To evaluate the performance of Reinforcement Learning and Genetic Algorithms, experiments have been conducted on the simple but effective environment of the cart pole.

Cart Pole is a classic control problem in reinforcement learning. The goal is to balance a pole on a cart that can move left or right.

The state space is four-dimensional, consisting of the cart position, cart velocity, pole angle, and pole angular velocity $[p, v, \alpha, \omega]$.

The action space is discrete, with two possible actions: move left or move right.

The reward is 1 for every time step the pole is balanced.

The goal is to balance the pole for as long as possible, with a limit of 500 actions.

The environment, called *CartPole-v1*, is implemented in Python using the Gymnasium library [towers'gymnasium'2023].

Below described implementations have been trained using the same environment and ensuring that in every iteration the starting point is the same between two methods, but different from the previous iteration, to ensure that the comparison can be evaluated without considering the stochastic nature of the training.

3.1. Reinforcement Learning

The reinforcement learning implementation is based on temporal difference learning [sutton1998temporal], in particular Q-learning. The implementation takes inspiration from the work of JackFurby [JackFurbyCartPole].

The *Q-table* is represented by the discretization of the continuous 4-dimensional state vector in 20 even intervals for every dimension of the vector leading to 160000 possible pairs of $\langle \text{state}, \text{action} \rangle$, considering the two possible actions.




Figure 4: Representation of the state discretization technique, considering an element of the 4D-state, s_i , the red dot is the real value of s_i , this is discretized to the nearest leftward discrete state.

Once an action is performed, the state selected is the first larger than the observed state.

The parameters used in the experiments are the following:

Learning rate α	0.1
Discount factor γ	1
Number of episodes n_{ep}	20000
Exploration rate ε	variable
Mutation Rate	0.05
Penalty factor PF	-375

Table 1: Parameters used in the RL implementation. The exploration rate ε starts with $\varepsilon(0) = 1$ and decays by $\varepsilon(t) = \varepsilon(t-1) - \frac{1}{\frac{n_{ep}}{2} - 1}$, every episode, stopping after $\frac{n_{ep}}{2}$ episodes.

3.2. Genetic Algorithms

Genotype

Since Genetic Algorithms can be very different depending on the genotype chosen to represent individuals, we have tried several different implementations of GA, varying the used genotype.

3.2.1. Encoding the sequence of actions

The first method used is a very naive implementation that can be applied to a very large variety of problems with GA: representing individuals with the vector of all actions they will perform in order. Thus, i -th character of the genotype of an individual j corresponds to the i -th action performed by the corresponding individual. In this approach, mutation is performed by switching an action in the genotype from left to right or from right to left with a probability given by the *Mutation rate* for every action i inside the genotype.

Since this particular genotype does not generalize well with the random initialization of the starting position of the pole. Due to its intrinsic dependency with the initial state, fixed starting conditions should be applied to effectively train in a meaningful way this genotype, by seeding the environment to always start in the same place, but this leads to a scarce ability of generalization, since the training

is valid just for a determined starting position of the pole.

All those considerations lead to the decision of evaluating other encodings for the final implementation.

3.2.2. Encoding a state-action table

The second encoding takes inspiration from Reinforcement Learning *Q-table*. In this approach, the focus is not to predict every action one by one but instead use GA to assign values to state-actions pairs then select the action that refers to the observed state.

The discretization technique is the same used in the reinforcement learning implementation and briefly described in figure ??.

Here, mutation is performed by swapping the action of a given state with a probability given by the *Mutation rate*.

Parameters

The GA parameters can be found in the following table :

Table 2: Parameters used in the GA implementation

Genotype	Q-table
Population Size	100
Generations	200
Selection	Fitness
Mutation Rate	0.005
Crossover	one-point
Elitism	2

4. Results

4.1. Training comparison




Figure 5: placeholder




Figure 6: placeholder

4.2. final model comparison




Figure 7: Difference between the two state action tables obtained using GA and RL. The black pixels represents the states where chosen action is the same, white pixels represent a different action's choice. The x-axis contains all the possible pairs of the first two elements of the state's 4D-vector, (s_1, s_2) , y-axis all the possible pairs of the last two elements (s_3, s_4) .




Figure 8: The plot shows the results obtained testing the two models on the same test set. The x-axis represents the number of the test set, the y-axis the number of steps the model was able to take before reaching the goal. The blue line represents the results obtained using the model trained with the GA, the red line the results obtained using the model trained with the RL.

5. Related Works

6. Discussion

Originally, when only trying the action-by-action approach in GA, observations led to one-sided results due to the dependency of the genotype on the initial state, leading to poor performances for GA implementation and no real generalization capacities. However, it was very interesting to see that the Genetic Algorithm performed much better when combining it with features from Reinforcement Learning - namely the *Q-table*.

The obtained results could be possible thanks to the *state discretization*, which allowed us to use the aforementioned tabular approach. Without the state discretization approach, different and more elaborate directions should have been considered. A discussed alternative could be exploring the path of function approximation, which is a very common approach in continuous state problems in RL, but no further experiments have been conducted in this direction due to the complex adaptations of this technique to the GA algorithm.

7. Conclusion

The conclusion should summarise your main results and main points from the discussion. A rule of thumb is to not present any new information (information not found in the results or discussion).

Acknowledgements

Acknowledgements (nb. takksigelse, nn. takkseingar) is not a requirement in a laboratory report. However, it is used in most scientific articles. It looks more professional and adds some “extra spice” to your report. Here is an example:

The authors would like to thank Dr. Ola Normann at the University of Oslo for assistance with the SIMS-analysis and Dr. Kari Normann at NTNU for fruitful discussions and support concerning melt spinning of silicon. This work was financially supported by the Norwegian research council and the Norwegian PhD Network on Nanotechnology for Microsystems.

Appendix A. Additional Information

You can use the appendix to include information that is relevant, but does not belong in the report. In most cases however, the appendix can be omitted and isn’t necessary.

Appendix A.1. Python code

If you used python code to process data, you can include the code (or a shorter version of it) in the appendix. Usually, however, it is better to hand in a separate file containing your code together with the report.ⁱ

Below is a simple example of some code used to calculate the values for the circuit in ??ⁱⁱ found in ??:

Listing 1: Example from external file

The code from ?? was displayed using a .py file. Since the lines are not numbered in this code example, you can copy and paste the code from the PDF into python without many issues (does however need to correct indents).

I would advice against using the `lstlisting` package to display code, as this introduces many unnecessary problems when trying to copy-paste the code.

Appendix B. Appendix footnotes

This template has a separate roman numeral footnote system for the appendix. You can chose to use this or normal footnotes in the appendix. Use

ⁱRule of thumb: short code goes in the appendix, long code goes in a separate file

ⁱⁱexample usage of the `varioref` package

the command `\appendixfootnote{text}`ⁱⁱⁱ to get a (lower case) roman numerical footnote. I added this footnote system because I thought it would be nice to have a separate footnote system for the appendix, since this section is in some ways separate from the rest of the document.

Appendix C. Boxes

This template also include two box environments to highlight text. I will showcase these in the two next subsections.

Appendix C.1. Info Box

The first environment is named `infobox` and is numbered, which allows for references to the box. You can also change the title of the box as well as the colours. To change the colour use the command `\SetInfoBoxBgColor{}` (changes background colour) and `\SetInfoBoxFrameColor{NTNU_blue}` (changes frame colour). The default colours are a light blue background and a darker blue frame. Here is an example:

Infobox	Appendix C-1
<p>Here is an infobox. You can also write math inside it:</p> $3x + 5y = 6z^2$	

Here is a reference to the infobox: ???. Notice that the structure of the infobox numbering is (section number)-(box number). The first infobox in section 2 thus has the reference 2-1.

Appendix C.2. Simple Box

The second environment is just a coloured box with no number or title. This can be used just to highlight text.

I also added a theorem environment `Sclaw` that may prove useful.

Scientific law 1 (Newton's 2. law).

$$\vec{F} = \frac{d\vec{p}}{dt}$$

You can reference the theorem environment: See ??. I also added a Norwegian version of the environment: `naturlov`. Let us change the colour of the next box to blue using `\SetSimpleBoxColor{bg_blue}`.

To create your own theorem environment, use the command `\newtheorem{}{}[]`.

You must use the `newtheorem` command before `\begin{document}` (the preamble). You can read more about the theorem environment in the [Overleaf documentation](https://www.overleaf.com/learn/latex/Theorems_and_proofs) using this link: https://www.overleaf.com/learn/latex/Theorems_and_proofs.

ⁱⁱⁱNote that there is **no** commands: `\appendixfootnotemark` and `\appendixfootnotetext`